

Generative AI MultiAgents



Learning Outcomes

Upon completion of this session, the learners should be able to:

- Explain the principles and benefits of multi-agent systems, including role specialization, task decomposition, communication protocols, and coordination strategies.
- Implement multi-agent workflows using advanced frameworks (e.g.CrewAI, LangGraph, AutoGen), incorporating shared memory management, orchestration, and failure recovery mechanisms.

Why we need to have Agents?

Complex decision-making:

Workflows involving nuanced judgment, exceptions, or context-sensitive decisions, for example refund approval in customer service workflows.

Difficult-to-maintain rules:

Systems that have become unwieldy due to extensive and intricate rulesets, making updates costly or error-prone, for example performing vendor security reviews.

Heavy reliance on unstructured data:

Scenarios that involve interpreting natural language, extracting meaning from documents, or interacting with users conversationally, for example processing a home insurance claim.

Before committing to building an agent, validate that your use case can meet these criteria clearly

What is a Generative AI MultiAgent?

Why Use Multiple Agents?


- **Single-agent limitation:** One LLM agent trying to solve everything often struggles with:
 - Context overload (too many instructions at once).
 - Weak specialization (generalist vs. expert).
 - Lack of modularity (hard to scale or debug).
- **Multi-agent benefits:**
 - **Role specialization** → different agents act as experts (e.g., planner, coder, reviewer).
 - **Parallelism** → tasks can be split among agents for faster execution.
 - **Modularity** → easier to upgrade or swap out one agent without redesigning the whole system.
 - **Human-like teamwork** → mirrors real-world collaboration.

Multi-Agent Planning & Task Decomposition

- **Hierarchical breakdown:**
 - Planner agent → breaks down the big problem into steps.
 - Worker agents → execute individual tasks.
 - Reviewer agent → checks correctness and quality.
- **Decomposition example: Build a website with product recommendations**
 - Planner: Defines steps → Data collection → Recommendation engine → Web design.
 - Data Agent: Retrieves product info.
 - ML Agent: Builds recommendation model.
 - UI Agent: Designs frontend.
 - Reviewer: Validates correctness.

Multiagent Pattern

- In the Router pattern, the router agent is using other agents as specialized tools, calling them one by one as needed.
- In the Parallel pattern, the initial agent is using multiple agents as tools simultaneously to speed up the process.
- In the Sequential and Circular Flow patterns, agents are used as tools in a predefined pipeline or loop.
- In the Dynamic pattern, agents are interacting more like a team, with each agent acting as both a user and a tool for other agents, depending on the situation.
- These patterns provide different ways to structure multi-agent interactions, allowing developers to choose the most appropriate approach based on the specific requirements of their application




Message Passing & Communication Protocols

- **Direct messaging:** Agent A sends structured text to Agent B.
- **Shared memory:** Agents post updates in a shared space.
- **Protocols:**
 - JSON-based structured messages
 - Tool invocation signals (like OpenAI's function calling)
 - Pub/Sub messaging (subscribe to updates on a task)
- **Key challenges:**
 - Ensuring clarity (avoid ambiguous prompts).
 - Preventing infinite loops (agents endlessly asking each other).
 - Controlling cost (token usage increases with more agents).

Frameworks & Tools

CrewAI

- Purpose: Organize **teams of agents** with defined roles and collaboration.
- Example roles: *Researcher, Analyst, Writer, Critic*.
- Supports **task delegation + team workflows**.
- Example use case: Generate a market report with multiple specialized agents.



Frameworks & Tools

LangGraph

- Model agent workflows as **state machines / graphs**.
- Nodes = agents or functions.
- Edges = transitions (who speaks next, what condition triggers a move).
- Good for **deterministic orchestration** (avoids chaos of “agents talking freely”).
- Example: Chatbot that can switch between *Sales Agent* → *Technical Agent* → *Escalation Agent* based on user intent.



Frameworks & Tools

AutoGen (Microsoft)


- Multi-agent conversation framework.
- Agents can:
 - Collaborate (e.g., Coder ↔ Reviewer loop).
 - Autonomously call tools.
 - Escalate to human-in-the-loop.
- Supports “**GroupChat**” pattern for open discussion among multiple agents.



Frameworks & Tools

LlamaIndex Agent

- Builds **data-aware agents** that can reason over structured + unstructured data.
- Uses **indexes** (vector, keyword, graph) for memory and context grounding.
- Supports **tool integration** (DB queries, APIs, Firecrawl, etc.).
- Works well as the **Knowledge Agent** inside multi-agent workflows.
- Example: Research Assistant agent that searches documents, summarizes results, and shares insights with a Planner or Reviewer agent.



Control & Coordination

Agent Orchestration

- **Centralized orchestration:** One “controller” agent assigns tasks.
- **Decentralized orchestration:** Agents coordinate via rules/messages, no single leader.
- Orchestration prevents chaos and ensures progress.

Managing Shared Memory or Knowledge

- Options:
 - Centralized memory store (vector DB, JSON file, blackboard).
 - Agent-specific memory (each agent has its own log/history).
 - Hybrid: Agents keep local context but sync key insights to global memory.
- Example: Planner maintains global plan, workers only store their progress.

Control & Coordination

Stopping Conditions & Failure Recovery

- Stopping conditions:
 - Goal achieved (final output validated).
 - Time/cost exceeded (budget control).
 - Deadlock detected (agents stop responding meaningfully).
- Failure recovery:
 - Retry mechanism with backoff.
 - Escalate to human (human-in-the-loop).
 - Replace failed agent with fallback.

References

- <https://www.voicespin.com/glossary/voice-activity-detection/>
- <https://www.youtube.com/watch?v=xWhI8RkRSGQ&t=60s>